

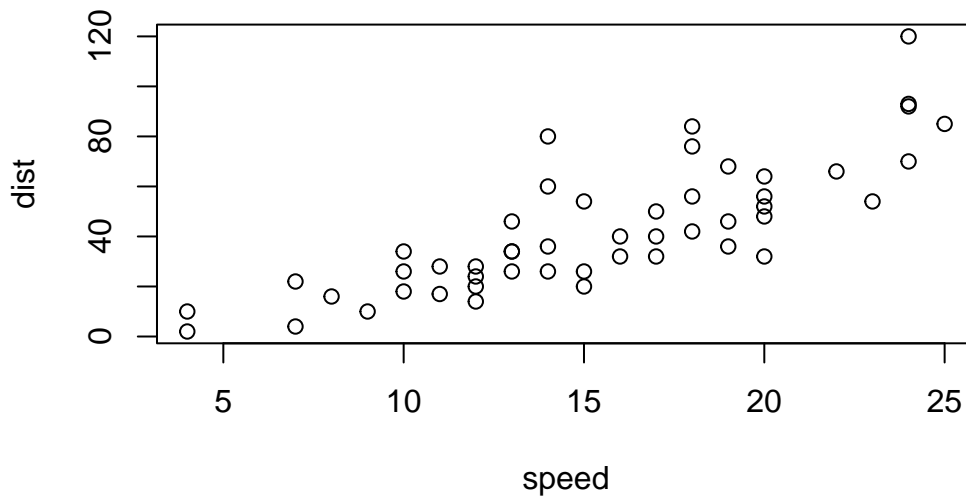
Class 5: Data Viz with ggplot

Ian Gurholt (PID A16767484)

Plotting in R

R has lots of ways to make plots and figures. This includes so-called **base** graphics and packages like **ggplot**

```
plot(cars)
```



This is a **base** R plot of the in-built `cars` dataset that has only two columns

```
head(cars)
```

	speed	dist
1	4	2
2	4	10
3	7	4
4	7	22
5	8	16
6	9	10

Q. How would we plot this data set with **ggplot2**?

All ggplot figures have at least 3 layers:

-data -aesthetics (how the data map to the plot) - geoms (how we draw the plot, lines, points, etc.)

Before I use any new package I need to download and install it with the `install.packages()` command.

I never use `install.packages()` within my quarto document otherwise I will install the package over and over and over again - which is silly!

Once a package is installed I can load it up with the `library()` function

Q1. For which phases is data visualization important in our scientific workflows?

All the above (Communication of Results, EDA, Detection of Outliers)

Q2. True or False? The ggplot2 package comes already installed with R?

FALSE

Q3. Which plot types are typically NOT used to compare distributions of numeric variables?

Network Graphs

Q4. Which statement about data visualization with ggplot2 is incorrect?

ggplot2 is the only way to create plots in R


Q5. Which geometric layer should be used to create scatter plots in ggplot2?

`geom_point()`

```
#install.packages(ggplot2)
library(ggplot2)
```

Warning: package 'ggplot2' was built under R version 4.3.3

```
ggplot(cars)+  
  aes(x=speed, y=dist)+  
  geom_point()
```



Key point: For simple plots (like the one above) ggplot is more verbose (without need to do more typing) but as plots get more complicated ggplot starts to be more clear and simple than base R plot()

Q6. In your own RStudio can you add a trend line layer to help show the relationship between the plot variables with the `geom_smooth()` function?

Refer to code for scatterplot

Q7. Argue with `geom_smooth()` to add a straight line from a linear model without the shaded standard error region?


Refer to code for scatterplot

Q8. Can you finish this plot by adding various label annotations with the `labs()` function and changing the plot look to a more conservative “black & white” theme by adding the `theme_bw()` function?

Refer to code for scatterplot

```
library(ggplot2)
ggplot(cars)+
  aes(speed,dist)+
  geom_point()+
  geom_smooth(se=FALSE, method=lm )+
  labs(title="Stopping distance of old cars",
        subtitle = "From the in-built cars dataset")+
  theme_bw()
```

`geom_smooth()` using formula = 'y ~ x'



```
url <- "https://bioboot.github.io/bimm143_S20/class-material/up_down_expression.txt"
genes <- read.delim(url)
head(genes)
```

	Gene	Condition1	Condition2	State
1	A4GNT	-3.6808610	-3.4401355	unchanging
2	AAAS	4.5479580	4.3864126	unchanging
3	AASDH	3.7190695	3.4787276	unchanging

```

4      AATF  5.0784720  5.0151916  unchanging
5      AATK  0.4711421  0.5598642  unchanging
6 AB015752.4 -3.6808610 -3.5921390  unchanging

```

Q9. Use the `nrow()` function to find out how many genes are in this dataset. What is your answer?

```
nrow(genes)
```

```
[1] 5196
```

There are 5196 genes in this dataset

Q10. Use the `colnames()` function and the `ncol()` function on the genes data frame to find out what the column names are (we will need these later) and how many columns there are. How many columns did you find?

```
colnames(genes)
```

```
[1] "Gene"      "Condition1" "Condition2" "State"
```

```
ncol(genes)
```

```
[1] 4
```

I found 4 columns in this dataset

Q11. Use the `table()` function on the State column of this data.frame to find out how many 'up' regulated genes there are. What is your answer?

```
table(genes$State)
```

```

down  unchanging      up
   72      4997      127

```

There are 127 upregulated genes

Q12. Using your values above and 2 significant figures. What fraction of total genes is up-regulated in this dataset?


```
round(table(genes$State)/nrow(genes) * 100, 2 )
```

down	unchanging	up
1.39	96.17	2.44

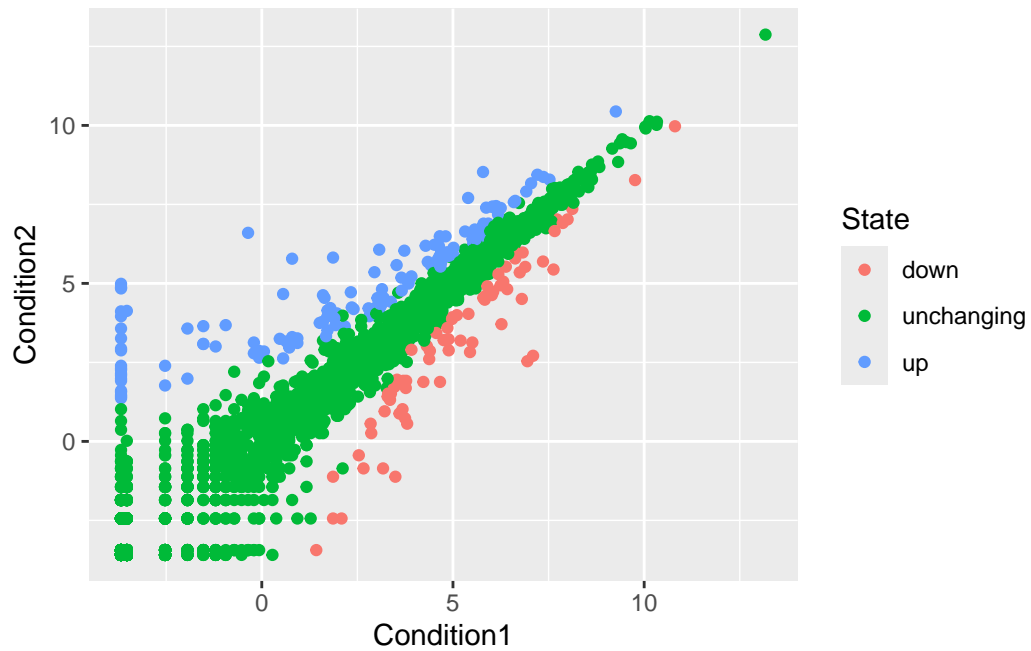
2.44% of genes are upregulated in this data set

Q13. Complete the code below to produce the following plot


```
ggplot(genes) +  
  aes(x=Condition1, y=Condition2) +  
  geom_point()
```



```
p <- ggplot(genes) +  
  aes(x=Condition1, y=Condition2, col=State) +  
  geom_point()  
p
```




```
p + scale_colour_manual( values=c("blue","gray","red") )
```



Q14. Nice, now add some plot annotations to the p object with the labs() function so your plot looks like the following:

```
p + labs(title = "Gene Expression Changes Upon Drug Treatment", x = "Control (no drug)", y = "Drug Treatment")
```



```
url <- "https://raw.githubusercontent.com/jennybc/gapminder/master/inst/extdata/gapminder.tsv"
gapminder <- read.delim(url)
```

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag


The following objects are masked from 'package:base':

intersect, setdiff, setequal, union



```
gapminder_2007 <- gapminder %>% filter(year==2007)
```

Q15. Complete the code below to produce a first basic scatter plot of this gapminder_2007 dataset:


```
library(ggplot2)
ggplot(gapminder_2007) +
  aes(x=gdpPerCap, y=lifeExp, color=continent, size=pop) +
  geom_point(alpha=0.5)
```



```
ggplot(gapminder_2007) +
  aes(x = gdpPerCap, y = lifeExp, color = pop) +
  geom_point(alpha=0.8)
```



```
ggplot(gapminder_2007) +  
  aes(x = gdpPerCap, y = lifeExp, size = pop) +  
  geom_point(alpha=0.5)+  
  scale_size_area(max_size = 10)
```

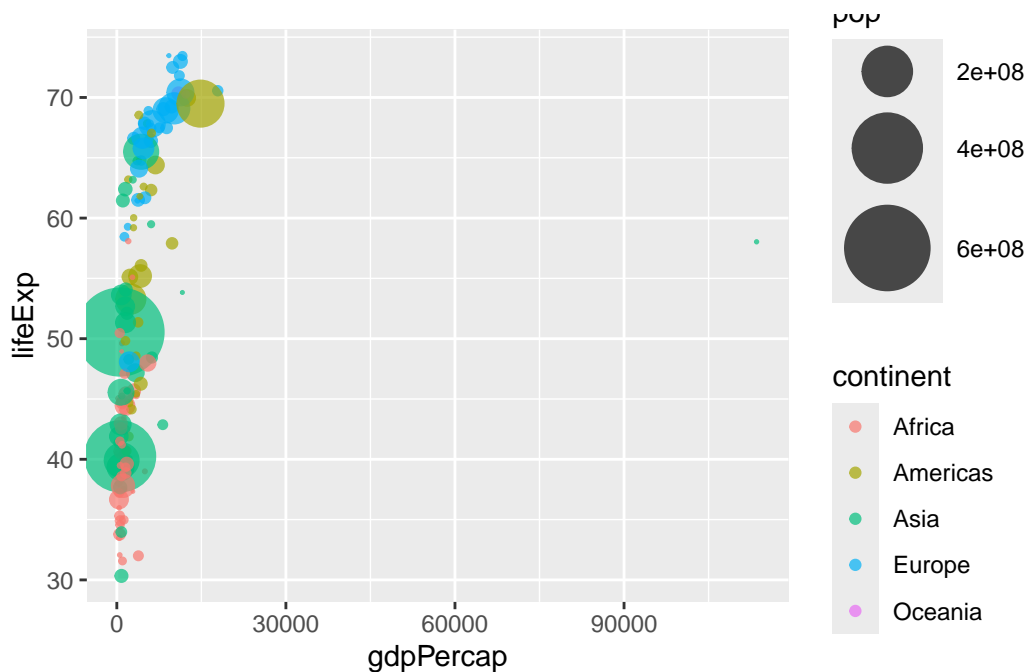


Q16. Can you adapt the code you have learned thus far to reproduce our gapminder scatter plot for the year 1957? What do you notice about this plot is it easy to compare with the one for 2007?

```
library(dplyr)

gapminder_1957 <- gapminder %>% filter(year==1957)

ggplot(gapminder_1957)+
  aes(x = gdpPercap, y = lifeExp, color=continent, size=pop)+
  geom_point(alpha=0.7)+
  scale_size_area(max_size = 15)
```




Without having both plots next to each other, or overlapping on the same graph, it is difficult to decipher the small differences between the two years as scrolling up and down to see each graph individually is inefficient and not very precise

Q17. Do the same steps above but include 1957 and 2007 in your input dataset for `ggplot()`. You should now include the layer `facet_wrap(~year)` to produce the following plot:

```
gapminder_1957 <- gapminder %>% filter(year==1957 | year==2007)

ggplot(gapminder_1957)+
  aes(x = gdpPercap, y = lifeExp, color=continent, size=pop)+
  geom_point(alpha=0.7)+
  scale_size_area(max_size = 10)+
  facet_wrap(~year)
```



. Q. Extract data for US in 1992

```
filter(gapminder, country=="United States",
       year==1992)
```

	country	continent	year	lifeExp	pop	gdpPerCap
1	United States	Americas	1992	76.09	256894189	32003.93

What was the population of Ireland in the last year we have data for?

```
filter(gapminder, country=="Ireland",
       year==2007)
```

	country	continent	year	lifeExp	pop	gdpPerCap
1	Ireland	Europe	2007	78.885	4109086	40676

Q. What countries in the data set had pop smaller than Ireland in 2007

```
gap07 <- filter(gapminder, year==2007)
ire_pop <- filter(gap07, country=="Ireland")["pop"]
ire_pop
```

```
pop
1 4109086
```

```
filter(gap07, pop<4109086)
```

	country	continent	year	lifeExp	pop	gdpPercap
1	Albania	Europe	2007	76.423	3600523	5937.0295
2	Bahrain	Asia	2007	75.635	708573	29796.0483
3	Botswana	Africa	2007	50.728	1639131	12569.8518
4	Comoros	Africa	2007	65.152	710960	986.1479
5	Congo, Rep.	Africa	2007	55.322	3800610	3632.5578
6	Djibouti	Africa	2007	54.791	496374	2082.4816
7	Equatorial Guinea	Africa	2007	51.579	551201	12154.0897
8	Gabon	Africa	2007	56.735	1454867	13206.4845
9	Gambia	Africa	2007	59.448	1688359	752.7497
10	Guinea-Bissau	Africa	2007	46.388	1472041	579.2317
11	Iceland	Europe	2007	81.757	301931	36180.7892
12	Jamaica	Americas	2007	72.567	2780132	7320.8803
13	Kuwait	Asia	2007	77.588	2505559	47306.9898
14	Lebanon	Asia	2007	71.993	3921278	10461.0587
15	Lesotho	Africa	2007	42.592	2012649	1569.3314
16	Liberia	Africa	2007	45.678	3193942	414.5073
17	Mauritania	Africa	2007	64.164	3270065	1803.1515
18	Mauritius	Africa	2007	72.801	1250882	10956.9911
19	Mongolia	Asia	2007	66.803	2874127	3095.7723
20	Montenegro	Europe	2007	74.543	684736	9253.8961
21	Namibia	Africa	2007	52.906	2055080	4811.0604
22	Oman	Asia	2007	75.640	3204897	22316.1929
23	Panama	Americas	2007	75.537	3242173	9809.1856
24	Puerto Rico	Americas	2007	78.746	3942491	19328.7090
25	Reunion	Africa	2007	76.442	798094	7670.1226
26	Sao Tome and Principe	Africa	2007	65.528	199579	1598.4351
27	Slovenia	Europe	2007	77.926	2009245	25768.2576
28	Swaziland	Africa	2007	39.613	1133066	4513.4806
29	Trinidad and Tobago	Americas	2007	69.819	1056608	18008.5092
30	Uruguay	Americas	2007	76.384	3447496	10611.4630
31	West Bank and Gaza	Asia	2007	73.422	4018332	3025.3498


** OPTIONAL BAR CHARTS**

```
gapminder_top5 <- gapminder %>%
  filter(year==2007) %>%
  arrange(desc(pop)) %>%
  top_n(5, pop)
```

```
gapminder_top5
```


	country	continent	year	lifeExp	pop	gdpPercap
1	China	Asia	2007	72.961	1318683096	4959.115
2	India	Asia	2007	64.698	1110396331	2452.210
3	United States	Americas	2007	78.242	301139947	42951.653
4	Indonesia	Asia	2007	70.650	223547000	3540.652
5	Brazil	Americas	2007	72.390	190010647	9065.801

```
ggplot(gapminder_top5) +
  geom_col(aes(x = country, y = pop))
```




Q18. Create a bar chart showing the life expectancy of the five biggest countries by population in 2007.


```
ggplot(gapminder_top5) +  
  geom_col(aes(x = country, y = lifeExp))
```



```
ggplot(gapminder_top5) +  
  geom_col(aes(x = country, y = pop, fill = continent))
```

```
ggplot(gapminder_top5) +  
  geom_col(aes(x = country, y = pop, fill = lifeExp))
```



Q19. Plot population size by country. Create a bar chart showing the population (in millions) of the five biggest countries by population in 2007.

```
ggplot(gapminder_top5) +  
  aes(x=reorder(country, -pop), y=pop, fill=country) +  
  geom_col(col="gray30") +  
  guides(fill="none")
```

