

# Class 7: Machine Learning

Ian Gurholt (PID: A16767484)

Today we are going to learn how to apply different machine learning methods, beginning with clustering:

The goal here is to find groups/clusters in your input data.


First I will make up some data with clear groups. For this I will use the 'rnorm()' function.

```
rnorm(10)
```


```
[1] -0.66578237  0.27938938  0.03948854 -0.41581878  0.07645015  0.79458540  
[7]  0.19938206 -0.42850706 -0.55673346  0.62889462
```

```
hist(rnorm(10000, mean=3))
```

**Histogram of rnorm(10000, mean = 3)**



```
x<-(c(rnorm(10000, -3), rnorm(10000,3)))
hist(x)
```




```
n<-30
x<-(c(rnorm(n, -3), rnorm(n,3)))
y <- rev(x)

z<-cbind(x,y)
head(z)
```

	x	y
[1,]	-4.193434	2.387783
[2,]	-3.294565	1.632687
[3,]	-4.069028	2.697352
[4,]	-3.292428	3.496142
[5,]	-2.339821	2.808686
[6,]	-2.277155	3.530184

```
plot(z)
```



Use the `kmeans()` function setting `k` to 2 and `nstart=20`

Inspect/print the results

Q. How many points are in each cluster? Q. What ‘component’ of your result object details - cluster size? - cluster assignment/membership? - cluster center?

Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
km<-kmeans(z, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-3.022063	3.142673
2	3.142673	-3.022063

Clustering vector:

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
```

```
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Within cluster sum of squares by cluster:

```
[1] 49.23019 49.23019
(between_SS / total_SS = 92.1 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Results in kmeansobject 'km'

```
attributes(km)
```

```
$names
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
$class
[1] "kmeans"
```

Cluster Size?

```
km$size
```

```
[1] 30 30
```

cluster assignment/membership?

```
km$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Cluster center?

```
km$centers
```

```
      x      y
1 -3.022063 3.142673
2  3.142673 -3.022063
```


Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

```
plot(z)
```




R will recycle the shorter color vector to be the same length as the longer (number of data points) in z


```
plot(z, col=c("red", "blue"))
```



```
plot(z, col= 2)
```




```
plot(z, col=km$cluster)
```



We can use the 'points()' function to add new points to an existing plot...like the cluster centers

```
plot(z, col=km$cluster)
points(km$centers, col= "blue", pch=15, cex=3)
```



Q. Can you run kmeans and ask for 4 clusters please and plot the results like we have done above

```
km4<-kmeans(z, centers=4)
km4
```

K-means clustering with 4 clusters of sizes 9, 4, 17, 30

Cluster means:

	x	y
1	-3.407327	2.266087
2	-4.864209	3.868938
3	-2.384653	3.435862
4	3.142673	-3.022063

Clustering vector:

```
[1] 1 1 1 3 3 3 1 3 1 3 3 2 2 1 3 3 3 3 3 1 1 3 2 3 3 3 3 1 3 2 4 4 4 4 4 4 4
[39] 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4
```

Within cluster sum of squares by cluster:


```
[1] 3.216211 2.977085 10.733295 49.230193
(between_SS / total_SS = 94.7 %)
```



Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"  
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
plot(z, col=km4$cluster)  
points(km4$centers, col= "blue", pch=15, cex=3)
```



## Heiarchical Clustering

Lets take our same made up data 'z' and see how hclust works

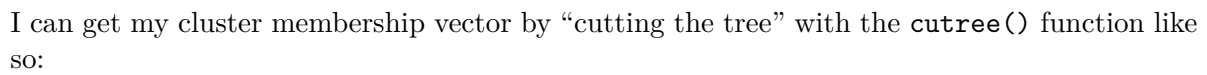
First we need a distance matrix of our data to be clustered.

```
d<- dist(z)  
hc<- hclust(d)  
hc
```

Call:


```
hclust(d = d)
```

```
plot(hc)
abline(h=8, col="red")
```



```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(z, col=grps)
```



## PCA of UK Food Data

Read data from UK dept of food consumption in different parts of the UK

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

Can use the `dim()` function in order to specify dimensions of a data frame

```
dim(x)
```

```
[1] 17  5
```

There are a total 17 rows and 5 columns in the original data frame

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.name=1)
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

I prefer to use the second method where we specified `row.names=1` meaning that we want to first column to be the names of the food and not the row numbers. This is a much quicker and efficient way as it is all condensed in a single function, however this may not be as robust as the first method as it only rids of numbered rows and if we wanted to get rid of the entire first column, given it was not just solely row numbers, we would then have to use the first method as it can remove the whole column and any content within it over and over again.

Maybe bar plots show some trends?


```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

Adjusting the `beside` function in the code from T to F will consequently make a vertically stacked bar plot instead of plots aligned next to each other

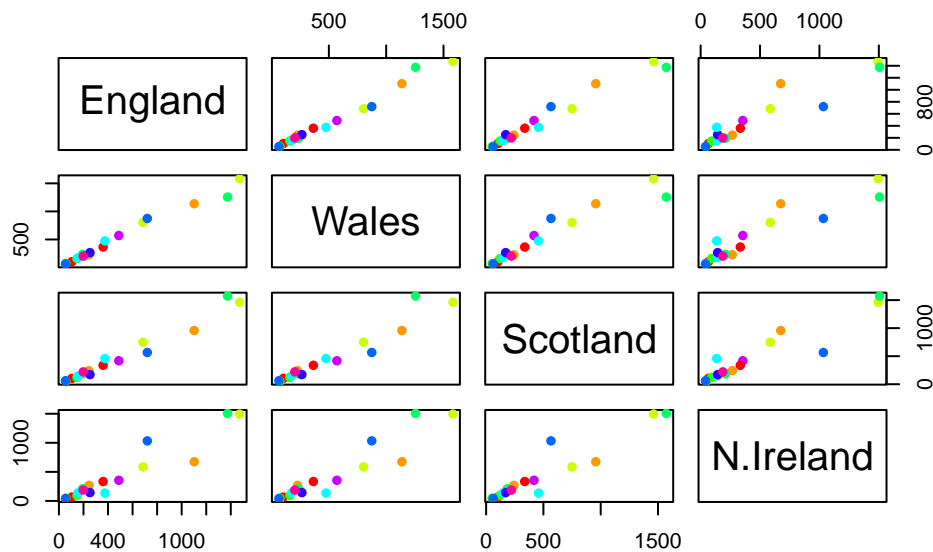
```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



No obvious trend with the barplot...

A so called Pairs Plot can be useful for small datasets like this one

```
pairs(x, col=rainbow(10), pch=16)
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

The code produces a pair plot by which in this case compares two countries and a given food variable to assess any differences between the two countries, if the points fall along the diagonal than one can assume that there are few differences between the two locations and similarities exist, however if the points fall above or below this diagonal, than differences do exist however the way these plots are represented makes it challenging to precisely understand the extent of these differences as many point lie very close to this diagonal and it is unclear whether one plot is more representative than the other.

Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Its hard to see structure and trends in even this small dataset. Many points appear to lie on the diagonal and some are just slightly off, however the extent of this variation is unclear. How will we ever do this when we have big datasets with 1000s of things we are measuring...

## PCA to the rescue

Let's see how PCA deals with this data set. So main function in base R to do PCA is called `prcomp()`

```
pca<- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what is inside this `pca` object that we created from running `prcomp()`

```
attributes(pca)
```

```
$names
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

```
$class
[1] "prcomp"
```


```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.


```
plot(pca$x[,1],pca$x[,2], col=c("black", "red", "blue","darkgreen"),pch=16, xlab="PC1 (67.4%")
text(pca$x[,1], pca$x[,2], colnames(x))
```





Q8. Customize your plot so that the colors of the country names match the colors in our UK and Ireland map and table at start of this document.

```
plot(pca$x[,1],pca$x[,2] ,pch=16, xlab="PC1 (67.4%)", ylab="PC2 (29%)", xlim=c(-270, 500))
text(pca$x[,1], pca$x[,2], colnames(x), col=c("orange", "red", "blue","darkgreen"))
```



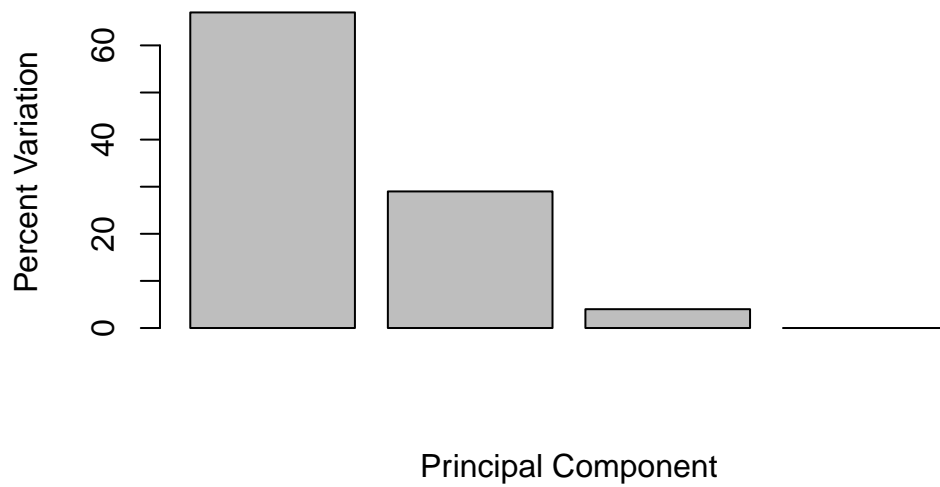
```
v <- round( pca$sdev^2/sum(pca$sdev^2) * 100 )
v
```

```
[1] 67 29 4 0
```

```
z <- summary(pca)
z$importance
```


	PC1	PC2	PC3	PC4
Standard deviation	324.15019	212.74780	73.87622	3.175833e-14
Proportion of Variance	0.67444	0.29052	0.03503	0.000000e+00
Cumulative Proportion	0.67444	0.96497	1.00000	1.000000e+00

```
barplot(v, xlab="Principal Component", ylab="Percent Variation")
```




##PCA Loading Scores

```
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Q9: Generate a similar 'loadings plot' for PC2. What two food groups feature prominently and what does PC2 mainly tell us about?

```
par(mar=c(10, 3, 0.35, 0))
barplot(pca$rotation[,2], las=2)
```



The two food groups that feature primarily in this graph is fresh potatoes and soft drinks and overall PC2 shows us a cleaner version of PC1 in which we can better decipher the skew of North Ireland to the right of the PCA being a result of soft drinks and the the skewing of the other countries to the left of the PCA being a result of the fresh potatoes negative value.

##PCA of RNA Seq Data

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

Q10: How many genes and samples are in this data set?

```
dim(rna.data)
```

```
[1] 100  10
```

Given that genes are rows and samples are columns, the `dim()` function shows that there are 100 genes and 10 samples